

Project 2 Final Report

Name - Vedant Kodagi

Email address - vkodagi@uchicago.edu

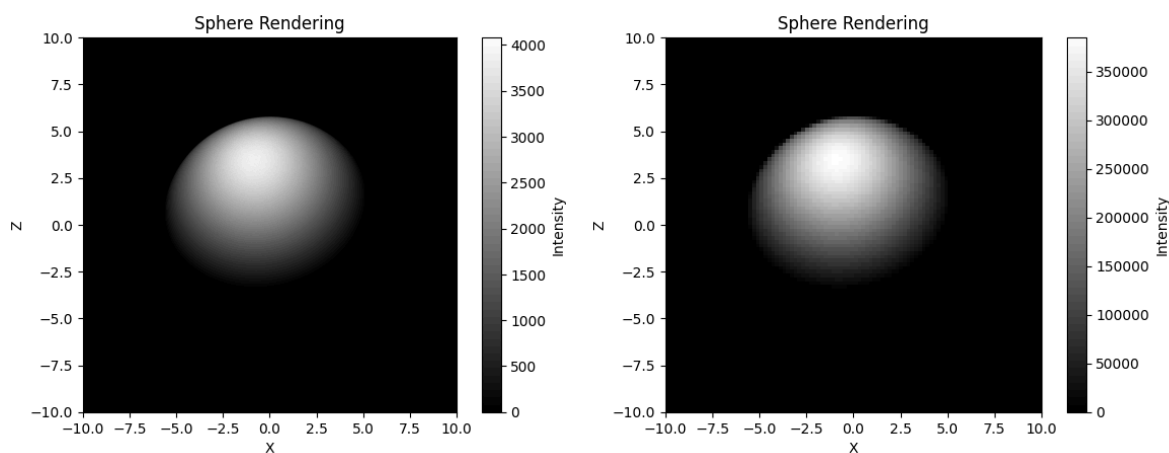
Professor - A. Seigel

Introduction

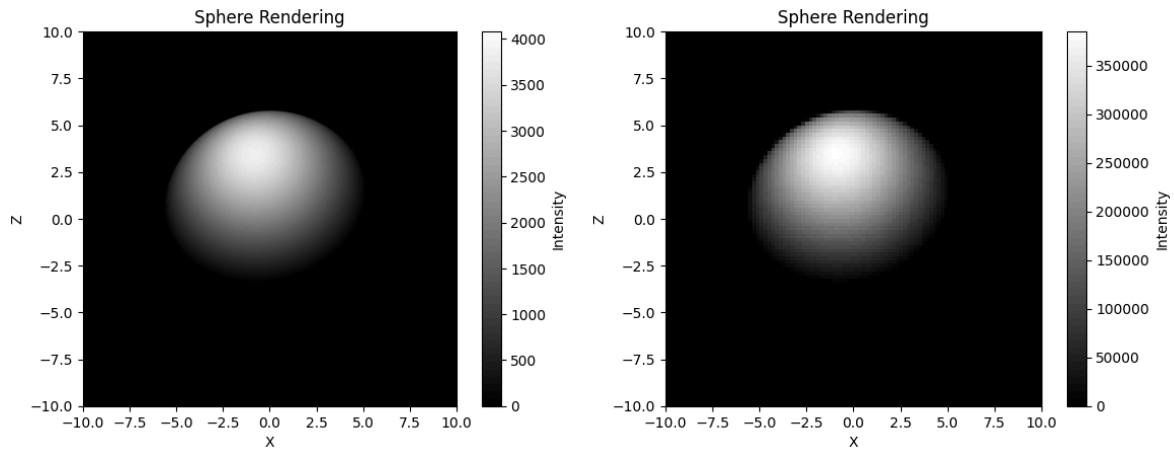
I did not alter my algorithmic approach as discussed in the last report however, I improved my kernel runtimes by nearly 4.7 times this week. I will discuss the details of this in the optimizations section.

Images by different processes-

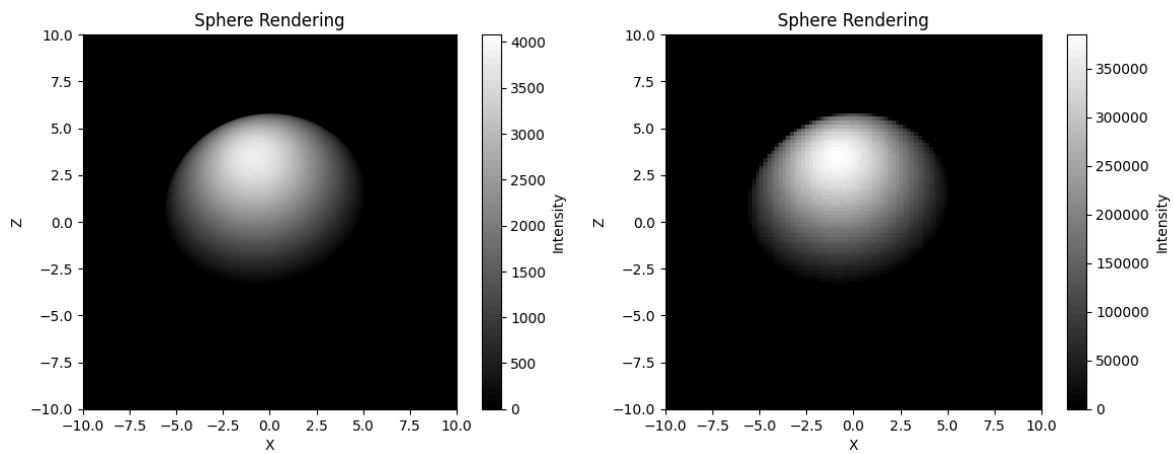
1. V100 images for 1000^2 and 100^2 grids.



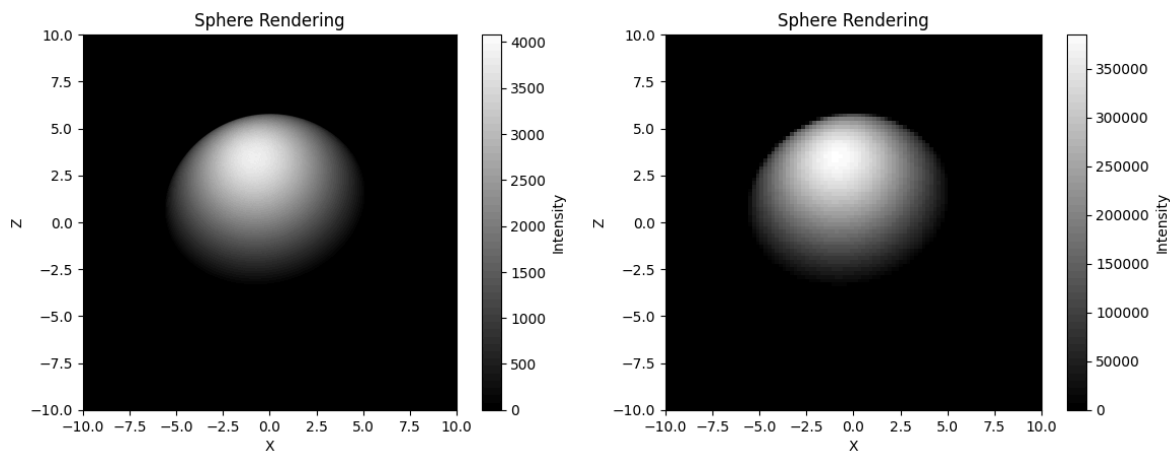
2. RTX6000 images for 1000^2 and 100^2 grids.



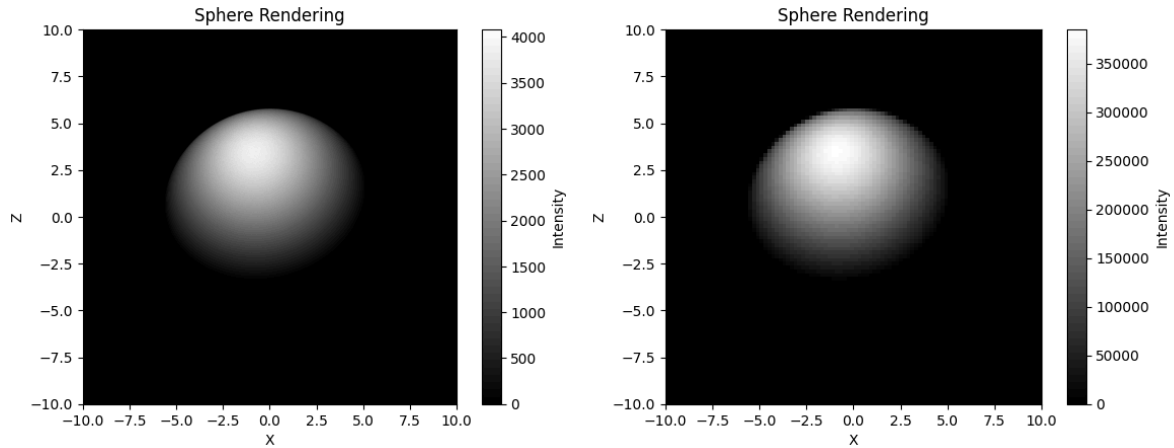
3. CPU Serial images for 1000^2 and 100^2 grids.



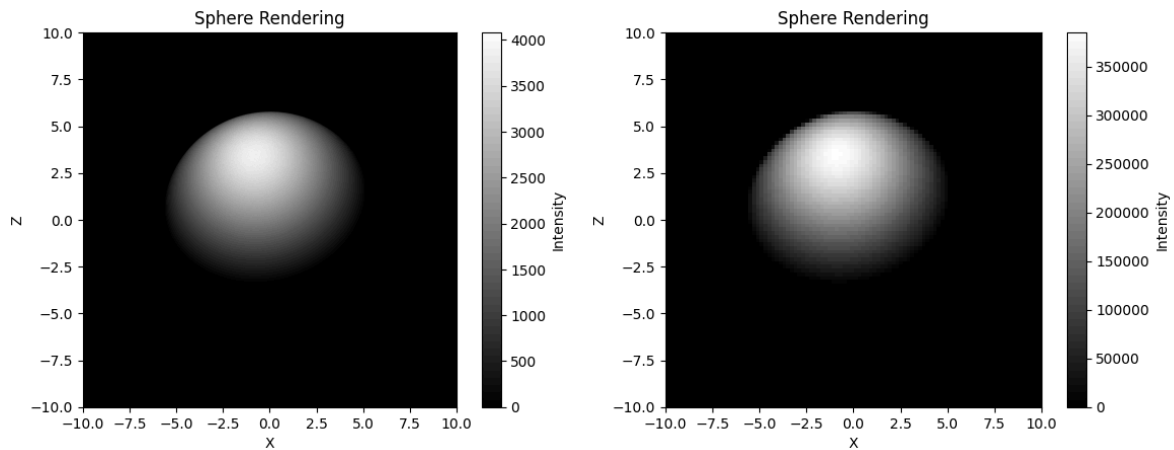
4. CPU OMP images for 1000^2 and 100^2 grids.



5. 2 V100 GPUs images for 1000^2 and 100^2 grids.



6. 4 V100 GPUs images for 1000^2 and 100^2 grids.



Performance

The fastest double-precision solution was 228.88 ms on a single A100 with a total wall time of 613 ms. Adding multiple GPUs, we got the expected speedups with 2 GPUs giving us x1.98 speedup and 4 GPUs giving us x3.85 speedup. However, it must be mentioned that there was a lot of variance in the runs. In 10 runs around 6 of them would give the runtimes mentioned in the table for the Multi-GPU cases. The number of blocks was always chosen to be equal to the number of streaming processors the GPU has. For example the V100 has 80 SMs, RTX6000 - 72 and A100 - 108. Threads per block were kind of a hit-and-trial method, I would check the average of say 5 runs and take whichever gave the fastest times.

Multi-GPU despite being considerably better than a single GPU, the total runtime was slower because I was using `MPI_Reduce()` to get the results from the GPUs which I think is the reason for the high total runtime. I am not sure whether just letting each of the GPUs write to a different file and then reducing the results would be faster.

Proc	Grid	Time (SP)	K Time (SP)	Time (DP)	K Time (DP)	Blk/TPB	Cores	Samples
A100	1000^2	503 ms	132.75 ms	613 ms	228.88 ms	108/512	-	14.83 Billion
A100	100^2	277 ms	132.75 ms	391 ms	228.84 ms	108/512	-	14.83 Billion
V100	1000^2	574 ms	178.01 ms	708 ms	312.04 ms	80/512	-	14.83 Billion
V100	100^2	341 ms	177.71 ms	475 ms	311.61 ms	80/512	-	14.83 Billion
RTX 6000	1000^2	851 ms	488.88 ms	5.05 s	4.68 s	72/256	-	14.91 Billion
RTX 6000	100^2	615 ms	487.66 ms	4.78 s	4.64 s	72/256	-	14.91 Billion
CPU Serial	1000^2	333.78 s	-	324.61 s	-	-	1	14.92 Billion
CPU Serial	100^2	333.55 s	-	317.52 s	-	-	1	14.92 Billion
CPU OMP	1000^2	26.54 s	-	25.37 s	-	-	16	14.92 Billion
CPU OMP	100^2	26.81 s	-	25.63 s	-	-	16	14.92 Billion
2 GPUs (V100)	1000^2	1.08 s	90.33 ms	1.20 s	157.43 ms	80/512	2	14.78 Billion
2 GPUs (V100)	100^2	0.81 s	90.31 ms	0.95 s	157.50 ms	80/512	2	14.78 Billion
4 GPUs (V100)	1000^2	1.75 s	47.23 ms	1.85 s	81.38 ms	80/512	4	14.87 Billion
4 GPUs (V100)	100^2	1.42 s	46.90 ms	1.50 s	83.10 ms	80/512	4	14.87 Billion

The performance tests were assumed with at least one billion rays (I mention at least because the way I divide work needed the number of rays to be divisible by the product of thread per block and the number of blocks, so I am calculating slightly more than a billion rays (a few thousand) for GPU problems), xorwow RNG in curand, with problem parameters set in Milestone 1 and 2. For both V100 and RTX 6000, I used -arch=sm_70. Anything other than this and the RTX 6000 would not launch the kernel. For A100 -arch=sm_80 was used.

Optimizations

The majority of my speedup came by simply changing the way I wrote my loops.

I went from this -

```
for(){  
    while(1){  
        // Get a ray that meets our conditions then break  
    }  
    // do other processing  
}
```

To-

```
while(count<number of rays){  
    if(a ray meets our conditions){  
        count++;  
        // do other processing  
    }  
}
```

This did not affect my serial runtimes however. But gives considerable speedup on a GPU. The next optimization I did was to reduce the number of times I call trigonometric functions. Using perf on my serial code I got the following output initially -

Samples: 1M of event 'cycles:u', Event count (approx.): 1307437720624

Overhead	Command	Shared Object	Symbol
39.03%	code	code	[.] main
34.47%	code	libm-2.28.so	[.] sincosf32x
11.87%	code	libc-2.28.so	[.] __random
8.85%	code	libc-2.28.so	[.] __random_r
3.51%	code	libc-2.28.so	[.] rand
1.18%	code	code	[.] rand@plt
0.57%	code	code	[.] sincos@plt
.			
.			

As you can see, the trigonometric functions were nearly 35 percent of the overhead. I was initially calling a sin function and a cos function for each sample. I reduced this to just one cos function and calculated the sin function value from the formula. This lead to the following perf report -

Samples: 1M of event 'cycles:u', Event count (approx.): 1152927769019

Overhead	Command	Shared Object	Symbol
48.36%	code	code	[.] main
13.55%	code	libc-2.28.so	[.] __random

10.18% code	libc-2.28.so	[.] __random_r
4.03% code	libc-2.28.so	[.] rand
2.52% code	libm-2.28.so	[.] __sqrt_finite
1.32% code	code	[.] rand@plt
0.71% code	libm-2.28.so	[.] sqrtf32x
0.70% code	libm-2.28.so	[.] 0x0000000000007ef2e
0.67% code	libm-2.28.so	[.] 0x0000000000007ee2f
0.66% code	libm-2.28.so	[.] 0x0000000000007ee21
0.66% code	libm-2.28.so	[.] 0x0000000000007ee13
0.65% code	code	[.] cos@plt

Reduced the overhead of trigonometric function calls to just 0.65 percent. This did lead to an increase in the square root function call, however the net overhead was reduced considerably.

I also changed from an array of pointers to a single contiguous array but I did not see any speedup with this. Maybe because we had added one more calculation to get the index of the array.